

Kubernatis (k8s) Basics

Kubernetes is an open-source platform designed to automate deploying, scaling, and operating application containers.

With Kubernetes,

Can quickly and efficiently respond to customer demand:

Deploy applications quickly and predictably.

Scale applications on the fly.

Roll out new features seamlessly.

Limit hardware usage to required resources only.

Goal is to foster an ecosystem of components and tools that relieve the burden of running applications in public and private clouds.

Kubernetes is

- Open source container cluster manager
- Developed by Google
- Used in google cloud
- Released in July 2015
- Goal is to automate deployment, scaling, maintain high availability
- Other Cluster management/Orchestration Tools:
 - Docker Swarm
 - Apache Mesos
- Owned by Seed Technology now

Components:

Nodes – minions
Pods
Labels
Selectors
Controllers
Services
Control Pane
API

Architecture:

Nodes can be virtual or physical hosts

Docker has to be installed on all of the nodes

Each minion will run ETCD (key pair management and communication service) for exchange of messages and reporting on cluster status

Pods

Pod consists of one or more containers, these containers are located on the same host.

They share the resources

Pods are assigned a unique IPs with in cluster

Pod management is done through the API

Labels:

Clients can attach key-value pairs to any object in the system

Grouping is done using labels and selectors

Controllers:

For management of cluster

Manage set of Pods depending on the desired state of the cluster

Service:

set of pods can work together , they can defined and implement a service ex: mysql or Apache

Installation and configuration

1. Install ntp on all the servers

yum install -y ntp

2. Start the service on all nodes

systemctl enable ntpd && systemctl start ntpd

3. Add ipaddress of each machine in /etc/hosts of all machines

```
[root@master ~]# cat /etc/hosts
```

```
10.0.0.81    master
```

10.0.0.144 minion1

10.0.0.61 minion2

4. Update file as below

vi /etc/yum.repos.d/virt7-docker-common-release.repo

```
[virt7-docker-common-release]
name=virt7-docker-common-release
baseurl=http://cbs.centos.org/repos/virt7-docker-common-release/x86_64/os/
gpgcheck=0
```

5. Enable the repo

***yum install -enablerepo=virt7-docker-common-release etcd
kubernetes docker -y***

Configuring Master:

=====

Add the information about the master controller

Configure 1 node as master

Configure 2 nodes as minions minion1 and minion2

1. Change the configuration file

vi /etc/kubernetes/config

Change the KUBE_MASTER as below

MASTER="--http://master:8080"

KUBE_ETCD_SERVERS="--etcd-servers=http://master:2379"

2. Changed the etcd configuration

vi /etc/etcd/etcd.conf

Change below:

In the [member] section

ETCD_LISTEN_CLIENT_URLS=<http://0.0.0.0:2379>

In the cluster section

ETCD_ADVERTISE_CLIENT_URLS=<http://0.0.0.0:2379>

3. Edit the API server

vi /etc/kubernetes/apiserver

```
KUBE_API_ADDRESS="--address=0.0.0.0"
Uncomment
KUBR_API_SEVER="--port=8080"
KUBELET_PORT="--kubelet-port=10250"
KUBE_SERVICE_ADDRESSES="--service-cluster-ip-range=10.254.0.0/16"
```

On master

```
systemctl enable etcd kube-apiserver kube-controller-manager
kube-scheduler
systemctl start etcd kube-apiserver kube-controller-manager
kube-scheduler
```

Configuring the minions:

=====

1. Change the configuration file
vi /etc/kubernetes/config

Change the KUBE_MASTER as below

```
MASTER="--http://master:8080"
KUBE_ETCD_SERVERS="--etcd-servers=http://master:2379"
```

2. Change the kubelet configuration
vi /etc/kubernetes/kubelet

```
KUBELET_ADDRESS="--address=0.0.0.0"
Uncomment kubelet port
KUBELET_PORT="--port=10250"
KUBELET_HOSTNAME="--hostname-override=minion1"
KUBELET_API_SERVER="--api-servers=http://master:8080"
Comment out KUBELET_POD_INFRA_CONTAINER
```

3. Start the kube services on all minions

```
systemctl enable kube-proxy kubelet docker
systemctl start kube-proxy kubelet docker
```

Perform same steps on all minions.

with this the configuration is complete

On master

```
[root@master ~]# kubectl get nodes
```

NAME	STATUS	AGE
minion1	Ready	4m
minion2	Ready	3m

```
[root@master ~]#
```

```
[root@master ~]# kubectl describe nodes
```

```
Name: minion1
Role:
Labels: beta.kubernetes.io/arch=amd64
        beta.kubernetes.io/os=linux
        kubernetes.io/hostname=minion1
Taints: <none>
CreationTimestamp: Sun, 15 Oct 2017 03:32:08 +0000
Phase:
Conditions:
```

Type	Status	LastHeartbeatTime
LastTransitionTime		Re
ason		Message
----	-----	-----
-----		--
----		-----
OutOfDisk	False	Sun, 15 Oct 2017
03:37:39 +0000	Sun, 15 Oct 2017	03:32:08 +0000
Ku		
beletHasSufficientDisk	kubelet has sufficient disk	
space available		
MemoryPressure	False	Sun, 15 Oct 2017
03:37:39 +0000	Sun, 15 Oct 2017	03:32:08 +0000
Ku		
beletHasSufficientMemory	kubelet has sufficient	
memory available		
DiskPressure	False	Sun, 15 Oct 2017
03:37:39 +0000	Sun, 15 Oct 2017	03:32:08 +0000
Ku		
beletHasNoDiskPressure	kubelet has no disk pressure	
Ready	True	Sun, 15 Oct 2017
03:37:39 +0000	Sun, 15 Oct 2017	03:32:18 +0000
Ku		
beletReady	kubelet is posting	
ready status		
Addresses:	10.0.0.144,10.0.0.144,minion1	

Capacity:

alpha.kubernetes.io/nvidia-gpu:	0
cpu:	2
memory:	500248Ki
pods:	110

Allocatable:

alpha.kubernetes.io/nvidia-gpu:	0
cpu:	2
memory:	500248Ki
pods:	110

System Info:

Machine ID:

980d6e7da5004216ade783778573fd3b

System UUID: D5F63352-C6CD-4EEF-

BB73-90368BD7677E

Boot ID: ccb89bc5-6da5-430c-

8da0-40899d370d84

Kernel Version: 3.10.0-

514.26.2.el7.x86_64

OS Image: CentOS Linux 7 (Core)

Operating System: linux

Architecture: amd64

Container Runtime Version: docker://1.12.6

Kubelet Version: v1.5.2

Kube-Proxy Version: v1.5.2

ExternalID: minion1

Non-terminated Pods: (0 in total)

Namespace		Name		CPU	
Requests	CPU Limits	Memory Requests	Memory Limits		
-----		----	-----		
-----	-----	-----	-----		

Allocated resources:

(Total limits may be over 100 percent, i.e., overcommitted.)

CPU Requests	CPU Limits	Memory Requests	Memory Limits
-----	-----	-----	-----

0 (0%)	0 (0%)	0 (0%)	0 (0%)

Events:

FirstSeen	LastSeen	Count	From	
SubObjectPath	Type	Reason		M
Message				
-----	-----	-----	----	
-----	-----	-----		-

11m	11m	1	{kube-proxy	
minion1}		Normal	Starting	
S				
tarting kube-proxy.				

5m	5m	1	{kubelet
minion1}		Normal	Starting

S

tarting kubelet.

5m	5m	1	{kubelet
minion1}		Warning	

ImageGCFailed u

nable to find data for container /

5m	5m	2	{kubelet
minion1}		Normal	

NodeHasSufficientD

isk Node minion1 status is now:

NodeHasSufficientDisk

5m	5m	2	{kubelet
minion1}		Normal	

NodeHasSufficientM

emory Node minion1 status is now:

NodeHasSufficientMemory

5m	5m	2	{kubelet
minion1}		Normal	

NodeHasNoDiskPress

ure Node minion1 status is now:

NodeHasNoDiskPressure

5m	5m	1	{kubelet
minion1}		Normal	

NodeReady N

ode minion1 status is now: NodeReady

Name: minion2
Role:
Labels: beta.kubernetes.io/arch=amd64
beta.kubernetes.io/os=linux
kubernetes.io/hostname=minion2
Taints: <none>
CreationTimestamp: Sun, 15 Oct 2017 03:33:10 +0000
Phase:
Conditions:

Type	Status	LastHeartbeatTime
OutOfDisk	False	Sun, 15 Oct 2017 03:37:41 +0000
MemoryPressure	False	Sun, 15 Oct 2017 03:37:41 +0000

LastTransitionTime	Reason	Message
OutOfDisk	Insufficient free disk space	kubelet has sufficient disk space available

LastTransitionTime	Reason	Message
MemoryPressure	Memory pressure	kubelet has sufficient disk space available

beletHasSufficientMemory kubelet has sufficient
memory available

 DiskPressure False Sun, 15 Oct 2017
03:37:41 +0000 Sun, 15 Oct 2017 03:33:10 +0000
Ku

beletHasNoDiskPressure kubelet has no disk pressure
 Ready True Sun, 15 Oct 2017
03:37:41 +0000 Sun, 15 Oct 2017 03:33:20 +0000
Ku

beletReady kubelet is posting
ready status
Addresses: 10.0.0.61,10.0.0.61,minion2

Capacity:
 alpha.kubernetes.io/nvidia-gpu: 0
 cpu: 2
 memory: 500248Ki
 pods: 110

Allocatable:
 alpha.kubernetes.io/nvidia-gpu: 0
 cpu: 2
 memory: 500248Ki
 pods: 110

System Info:
 Machine ID:
980d6e7da5004216ade783778573fd3b

```

System UUID:                D5CE9FAF-C19F-43AF-
A2B8-8D2D07799368
Boot ID:                    7f676c28-3494-472c-
a900-e57bad33f042
Kernel Version:            3.10.0-
514.26.2.el7.x86_64
OS Image:                  CentOS Linux 7 (Core)
Operating System:          linux
Architecture:              amd64
Container Runtime Version:  docker://1.12.6
Kubelet Version:           v1.5.2
Kube-Proxy Version:        v1.5.2
ExternalID:                minion2
Non-terminated Pods:       (0 in total)

```

Namespace	Name	CPU
Requests	CPU Limits	Memory Requests Memory
Limits		
-----	----	-----
-----	-----	-----

Allocated resources:

(Total limits may be over 100 percent, i.e.,
overcommitted.)

CPU Requests	CPU Limits	Memory Requests	Memory
Limits			
-----	-----	-----	-----

0 (0%)	0 (0%)	0 (0%)	0 (0%)
Events:			
FirstSeen	LastSeen	Count	From
SubObjectPath	Type	Reason	M
Message			
-----	-----	-----	-----
-----	-----	-----	-

4m	4m	1	{kube-proxy
minion2}		Normal	Starting
S			
tarting kube-proxy.			
4m	4m	1	{kubelet
minion2}		Normal	Starting
S			
tarting kubelet.			
4m	4m	1	{kubelet
minion2}		Warning	
ImageGCFailed	u		
nable to find data for container /			
4m	4m	2	{kubelet
minion2}		Normal	
NodeHasSufficientD			
isk	Node minion2 status is now:		
NodeHasSufficientDisk			

```
4m          4m          2      {kubelet
minion2}          Normal
```

NodeHasSufficientM

emory Node minion2 status is now:

NodeHasSufficientMemory

```
4m          4m          2      {kubelet
minion2}          Normal
```

NodeHasNoDiskPress

ure Node minion2 status is now:

NodeHasNoDiskPressure

```
4m          4m          1      {kubelet
minion2}          Normal
```

NodeReady N

ode minion2 status is now: NodeReady

To get information about nodes:

kubect! get nodes

Creating POD using yaml file:

```
[root@master test]# cat nginx.yaml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: nginx
```

```
spec:
```

```
containers:
- name: nginx
  image: nginx:1.7.9
  ports:
  - containerPort: 80
```

To create POD:

```
kubectl create -f ./nginx.yaml
```

To list PODs:

```
Kubectl get pods
```

Creation of POD may fail because of a bug

Work around for bug in kubernetes:

=====

<https://github.com/kubernetes/kubernetes/issues/11355#issuecomment-127378691>

Restart all the process once a key is generated

Label is to identity the POD in thousands of PODS , it is like a tag

It is a key value pair

```
[root@master test]# cat nginx1.yaml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
  name: nginx1
```

```
  labels:
```

```
    app: nginx1
spec:
  containers:
  - name: nginx
    image: nginx:1.7.9
    ports:
    - containerPort: 80
```

kubectl create -f nginx1.yaml

Creates pod nginx1

kubectl get-pods -l app=nginx

Description of pods

```
[root@master test]# kubectl get pods -l app=nginx
```

NAME	READY	STATUS	RESTARTS	AGE
nginx3	1/1	Running	0	18s

```
[root@master test]# kubectl describe pods -l app=nginx
```

```
Name:          nginx3
Namespace:     default
Node:          minion2/10.0.0.40
Start Time:    Sat, 21 Oct 2017 02:27:45 +0000
Labels:        app=nginx
Status:        Running
IP:            172.17.0.4
Controllers:   <none>
Containers:
```


nginx1:

Container ID:

docker://6be47b2532269229d7fcc6b1a878b15ed3dd1e4019e489
4b8c375044c3934171

Image: nginx:1.7.9

Image ID: docker-

pullable://docker.io/nginx@sha256:e3456c851a152494c3e4f
f5fcc26f240206abac0c9d794aff
b40e0714846c451

Port: 80/TCP

State: Running

Started: Sat, 21 Oct 2017 02:27:48 +0000

Ready: True

Restart Count: 0

Volume Mounts:

/var/run/secrets/kubernetes.io/serviceaccount
from default-token-jglff (ro)

Environment Variables: <none>

Conditions:

Type	Status
Initialized	True
Ready	True
PodScheduled	True

Volumes:

default-token-jglff:

Type: Secret (a volume populated by a Secret)

```

    SecretName: default-token-jglff
QoS Class:      BestEffort
Tolerations:    <none>
Events:
  FirstSeen      LastSeen        Count   From
SubObjectPath               Type            Reason      M
essage
-----
-----
-----
1m              1m              1       {default-
scheduler }      Normal
ScheduledS
uccessfully assigned nginx3 to minion2
1m              1m              2       {kubelet
minion2}          Warning
MissingClu
sterDNS kubelet does not have ClusterDNS IP configured
and cannot create Pod using "ClusterFirst" policy.
Falling
back to DNSDefault policy.
1m              1m              1       {kubelet
minion2}          spec.containers{nginx1} Normal
Pulled    C
ontainer image "nginx:1.7.9" already present on machine

```

```
1m          1m          1          {kubelet
minion2}    spec.containers{nginx1} Normal
```

Created C

reated container with docker id 6be47b253226;

Security:[seccomp=unconfined]

```
1m          1m          1          {kubelet
minion2}    spec.containers{nginx1} Normal
```

Started S

tarted container with docker id 6be47b253226

Deployments:

For production:

For deployment, we need to use latest API which is in
Extension

```
[root@master test]# cat nginx-deploy-production.yaml
```

```
apiVersion: extensions/v1beta1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: nginx-deploy
```

```
spec:
```

```
  replicas: 3
```

```
template:
  metadata:
    labels:
      app: nginx-deploy
  spec:
    containers:
      - name: nginx-deploy
        image: nginx:1.7.9
        ports:
          - containerPort: 80
```

With deployment we can do rolling update to PODs,

Update nginx version from 1.7.9 to 1.8

Kubectrl apply nginx-deploy.yaml read yaml file and update to 1.8

Update production env to new level

Kind:

Pod to create a POD definition

Deployment For creating production like environment and perform update

Replicationcontroller for high availability

Service Acts as a load balancer to Kubernetes Cluster.

Kubernetes Pods are mortal. They are born and when they die, they are not resurrected. ReplicationControllers in particular create and destroy Pods dynamically (e.g. when scaling up or down or when doing rolling updates). While each Pod gets its own IP address, even those IP addresses cannot be relied upon to be stable over time.

This leads to a problem: if some set of Pods (let's call them backends) provides functionality to other Pods (let's call them frontends) inside the Kubernetes cluster,

how do those frontends find out and keep track of which backends are in that set?

Solution is services

A Kubernetes Service is an abstraction which defines a logical set of Pods and a policy by which to access them - sometimes called a micro-service. The set of Pods targeted by a Service is (usually) determined by a Label Selector (see below for why you might want a Service without a selector).

As an example, consider an image-processing backend which is running with 3 replicas. Those replicas are fungible - frontends do not care which backend they use. While the actual Pods that compose the backend set may change, the frontend clients should not need to be aware of that or keep track of the list of backends themselves. The Service abstraction enables this decoupling.

For Kubernetes-native applications, Kubernetes offers a simple Endpoints API that is updated whenever the set of Pods in a Service changes. For non-native applications, Kubernetes offers a virtual-IP-based bridge to Services which redirects to the backend Pods.

Documentation :

<https://kubernetes.io/docs/concepts/services-networking/service/#type-nodeport>

Examples:

Example:

```
$ vi nginx_pod.yaml
apiVersion: v1
kind: ReplicationController
metadata:
```

```
  name: nginx
spec:
  replicas: 2
  selector:
    app: nginx
  template:
    metadata:
      name: nginx
      labels:
        app: nginx
    spec:
      containers:
      - name: nginx
        image: nginx
        ports:
        - containerPort: 80
```

```
apiVersion: v1
kind: Service
metadata:
  name: api
  namespace: wardle
spec:
  ports:
  - port: 443
    protocol: TCP
    targetPort: 443
  selector:
```

apiserver: "true"

```
[root@minion1 ~]# docker ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS
		PORTS		NAMES
ae3fc2bb7b68	nginx	"nginx -g 'daemon off'"	35 minutes ago	Up 35 minutes
k8s_nginx.a6022f15_nginx-gr8ck_default_0535328e-b73e-11e7-8a52-080027ce8866_70008e87				
2d3378f028ca	gcr.io/google_containers/pause-amd64:3.0	"/pause"	35 minutes ago	Up 35 minutes
k8s_POD.b2390301_nginx-gr8ck_default_0535328e-b73e-11e7-8a52-080027ce8866_0cb243ea				

```
[root@minion1 ~]#
```

A pause container is created by default on minion

Why?

In Kubernetes, each pod has an IP and within a pod there exists a so called infrastructure container, which is the first container that the Kubelet instantiates and it acquires the pod's IP and sets up the network namespace. All the other containers in the pod then join the infra container's network and IPC namespace. The infra container has network bridge mode enabled and all the other containers in the pod share its namespace via container mode. The initial process that runs in the infra container does effectively nothing since its sole purpose is to act as the home for the namespaces.

Nginx Server Deployment using Kubernetes

=====

1. Create pod yaml file

vi nginx_pod.yaml

apiVersion: v1

kind: ReplicationController

metadata:

name: nginx

spec:

replicas: 2

selector:

app: nginx

template:

metadata:

name: nginx

labels:

app: nginx

spec:

containers:

- name: nginx

image: nginx

ports:

- containerPort: 80

2. Create pod

kubectrl create -f nginx_pod.yaml

3. Deploy the nginx service using yaml file in order to expose the nginx pod on the host port "82"

\$ vi nginx_service.yaml

apiVersion: v1

kind: Service

metadata:

labels:

name: nginxservice

name: nginxservice

spec:

ports:

The port that this service should serve on.

- port: 82

Label keys and values that must match in order to receive traffic for this service.

selector:

app: nginx

type: LoadBalancer

4. Create the nginx service using kubectl

kubectl create -f nginx_service.yaml

services/nginxservice

The nginx service can be listed as follow

kubectl get services

NAME	LABELS	SELECTOR	IP(S)	PORT(S)
kubernetes	component=apiserver,provider=kubernetes		<none>	192.168.3.1 443/TCP
nginxservice	name=nginxservice	app=nginx		192.168.3.43 82/TCP

